





Image Primitives and Correspondence

Image



Brightness values



$$I(x,y)$$

Image Features

Local, meaningful, detectable parts of the image.

- Edge detection
- Line detection
- Corner detection

Motivation

- Information content high
- Invariant to change of view point, illumination
- Reduces computational burden
- Uniqueness
- Can be tuned to a task at hand

Filtering and Image Features



Given a noisy image

How do we reduce noise ?

How do we find useful features ?

Today:

- Filtering
- Point-wise operations
- Edge detection



Moving average

- Let's replace each pixel with a *weighted* average of its neighborhood
- The weights are called the *filter kernel*
- What are the weights for the average of a 3x3 neighborhood?

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1


“box filter”

Defining convolution

- Let f be the image and g be the kernel. The output of convolving f with g is denoted $f * g$.

$$(f * g)[m, n] = \sum_{k, l} f[m - k, n - l] g[k, l]$$


Convention:
kernel is “flipped”



- MATLAB functions: [conv2](#), [filter2](#), [imfilter](#)

Details

- What is the size of the output?
- MATLAB: `filter2(g, f, shape)`
 - *shape* = 'full': output size is sum of sizes of *f* and *g*
 - *shape* = 'same': output size is same as *f*
 - *shape* = 'valid': output size is difference of sizes of *f* and *g*



Averaging filter 1-D example

$$g[x] = \sum_{k=-\infty}^{\infty} f[k]h[x - k]$$

$$f[x] = [...0, 0, 2, -2, 2, 0, 0, ...] \quad h[x] = \frac{1}{3}[1, 1, 1]$$

$$h[-1] = \frac{1}{3}, h[0] = \frac{1}{3}, h[1] = \frac{1}{3} \quad \text{and 0 everywhere else}$$

$$f[-1] = -2, f[0] = 2, f[1] = -2$$

Box filter

$$g[x] = \sum_{k=-1}^1 f[k]h[x - k]$$

Ex. cont.





$$g[-1] = f[-1]h[-1 - 1] + f[0]h[-1] + f[1]h[0]$$

$$g[0] = f[-1]h[-1] + f[0]h[0] + f[1]h[1]$$

Averaging filter center pixel weighted more

$$h[x] = [0.25, 0.5, 0.25]$$

Averaging filter



Convolution in 2D

$$g[x, y] = \sum_{k=-\frac{w}{2}}^{\frac{w}{2}} \sum_{l=-\frac{w}{2}}^{\frac{w}{2}} f[k, l] h[x - k, y - l]$$

f

10	11	10	0	0	1
9	10	11	1	0	1
10	9	10	0	2	1
11	10	9	10	9	11
9	10	11	9	99	11
10	9	9	11	10	10

h

1	1	1
1	1	1
1	1	1

g

X	X	X	X	X	X
X	10				X
X					X
X					X
X					X
X	X	X	X	X	X

$$1/9.(10 \times 1 + 11 \times 1 + 10 \times 1 + 9 \times 1 + 10 \times 1 + 11 \times 1 + 10 \times 1 + 9 \times 1 + 10 \times 1) = 1/9.(90) = 10$$

Example:

I

10	11	10	0	0	1
9	10	11	1	0	1
10	9	10	0	2	1
11	10	9	10	9	11
9	10	11	9	99	11
10	9	9	11	10	10

F

1	1	1
1	1	1
1	1	1

O

X	X	X	X	X	X
X	10	7	4	1	X
X					X
X					X
X					X
X	X	X	X	X	X

$$1/9.(10x1 + 0x1 + 0x1 + 11x1 + 1x1 + 0x1 + 10x1 + 0x1 + 2x1) = 1/9.(34) = 3.7778$$

Example:

I

10	11	10	0	0	1
9	10	11	1	0	1
10	9	10	0	2	1
11	10	9	10	9	11
9	10	11	9	99	11
10	9	9	11	10	10

O

X	X	X	X	X	X
X	10	7	4	1	X
X					X
X					X
X				20	X
X	X	X	X	X	X

F

1	1	1
1	1	1
1	1	1

$$1/9.(10 \times 1 + 9 \times 1 + 11 \times 1 + 9 \times 1 + 99 \times 1 + 11 \times 1 + 11 \times 1 + 10 \times 1 + 10 \times 1) = 1/9.(180) = 20$$

Example:

I

10	11	10	0	0	1
9	10	11	1	0	1
10	9	10	0	2	1
11	10	9	10	9	11
9	10	11	9	99	11
10	9	9	11	10	10

O

X	X	X	X	X	X
X	10	7	4	1	X
X					X
X			18		X
X				20	X
X	X	X	X	X	X

F


1	1	1
1	1	1
1	1	1

$$1/9.(10 \times 1 + 0 \times 1 + 2 \times 1 + 9 \times 1 + 10 \times 1 + 9 \times 1 + 11 \times 1 + 9 \times 1 + 99 \times 1) = 1/9.(159) = 17.6667$$

How big should the mask be?

- The bigger the mask,
 - more neighbors contribute.
 - smaller noise variance of the output.
 - bigger noise spread.
 - more blurring.
 - more expensive to compute.
 - In Matlab function **conv**, **conv2**

Example: Smoothing by Averaging




Gaussian Filter

- A particular case of averaging
 - The coefficients are samples of a 1D Gaussian.
 - Gives more weight at the central pixel and less weights to the neighbors.
 - The further away the neighbors, the smaller the weight.

$$g(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{\frac{-x^2}{2\sigma^2}},$$

Sample from the continuous Gaussian

Smoothing with a Gaussian




How big should the mask be?

- The std. dev of the Gaussian σ determines the amount of smoothing.
- The samples should adequately represent a Gaussian
- For a 98.76% of the area, we need

$$m = 5\sigma$$

$$5 \cdot (1/\sigma) \leq 2\pi \Rightarrow \sigma \geq 0.796, m \geq 5$$

5-tap filter



$$g[x] = [0.136, 0.6065, 1.00, 0.606, 0.136]$$

Gaussian filters

- Remove “high-frequency” components from the image (low-pass filter)
- Convolution with self is another Gaussian
 - So can smooth with small- σ kernel, repeat, and get same result as larger- σ kernel would have
 - Convoluting two times with Gaussian kernel with std. dev. σ
is same as convoluting once with kernel with std. dev. $\sigma\sqrt{2}$
- *Separable* kernel
 - Factors into product of two 1D Gaussians

Separability of the Gaussian filter

$$\begin{aligned} G_{\sigma}(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

In this case, the two functions are the (identical) 1D Gaussian

Separability example

2D convolution
(center location only)

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix}$$

The filter factors
into a product of 1D
filters:

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \times \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

Perform convolution
along rows:

$$\begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix} = \begin{bmatrix} & 11 & \\ & 18 & \\ & 18 & \end{bmatrix}$$

Followed by convolution
along the remaining column:



Image Smoothing

- Convolution with a 2D Gaussian filter

$$\tilde{I}(x, y) = I(x, y) * g(x, y) = I(x, y) * g(x) * g(y)$$

- Gaussian filter is separable, convolution can be accomplished as two 1-D convolutions

$$\tilde{I}[x, y] = I[x, y] * g[x, y] = \sum_{k=-\frac{w}{2}}^{\frac{w}{2}} \sum_{l=-\frac{w}{2}}^{\frac{w}{2}} I[k, l] g[x - k] g[y - l]$$




How big should the mask be?


- The bigger the mask,
 - more neighbors contribute.
 - smaller noise variance of the output.
 - bigger noise spread.
 - more blurring.
 - more expensive to compute.

Edges


- They happen at places where the image values exhibit sharp variation



Edge detection (1D)




Edge= sharp variation




Large first derivative

Digital Approximation of 1st derivatives

$$\frac{df(x)}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x + \Delta x) - f(x)}{\Delta x}$$



$$\frac{df(x)}{dx} \cong \frac{f(x+1) - f(x-1)}{2}$$



Convolve with:

-1	0	1
----	---	---

Edge Detection (2D)

Vertical Edges:

Convolve with:


-1	0	1
----	---	---

Horizontal Edges:

Convolve with:

-1
0
1

Noise cleaning and Edge Detection



- we need to also deal with noise
- Combine Linear Filters
- Instead of smoothing, followed by derivative computation
- Convolve with derivative of the smoothing filter

Noise Smoothing & Edge Detection

Convolve with:

-1	0	1
-1	0	1
-1	0	1

Vertical Edge Detection

Noise Smoothing

This mask is called the (vertical) Prewitt Edge Detector

Outer product of box filter $[1 \ 1 \ 1]^T$ and $[-1 \ 0 \ 1]$

Noise Smoothing & Edge Detection

Convolve with:


-1	-1	-1
0	0	0
1	1	1

Noise Smoothing


Horizontal Edge Detection

This mask is called the (horizontal) Prewitt Edge Detector

Gaussian and its derivative




$$g(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{\frac{-x^2}{2\sigma^2}}, \quad g'(x) = -\frac{x}{\sigma^2\sqrt{2\pi}\sigma} e^{\frac{-x^2}{2\sigma^2}}.$$




Vertical edges

$$I_x(x, y) = \frac{\partial I}{\partial x}$$



First derivative - one column



Horizontal edges

$$I_y(x, y) = \frac{\partial I}{\partial y}$$



Gradient orientation



- Image Gradient $\nabla I = \left[\frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right]$



- Gradient Magnitude

$$m = \sqrt{\left(\frac{\partial I}{\partial x} \right)^2 + \left(\frac{\partial I}{\partial y} \right)^2}$$

- Gradient Orientation

$$\theta = \tan^{-1} \left(\frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right)$$

Canny Edge Detector



- Edge detection involves 3 steps:
 - Noise smoothing
 - Edge enhancement
 - Edge localization
- J. Canny formalized these steps to design an *optimal* edge detector
- How to go from derivatives to edges ?

Horizontal edges

Edge Detection



original image



gradient magnitude


Canny edge detector

- Compute image derivatives
- if gradient magnitude $> \tau$ and the value is a local maximum along gradient direction – pixel is an edge candidate

Algorithm Canny Edge detector

- The input is image I ; G is a zero mean Gaussian filter (std = σ)
 - 1. $J = I * G$ (smoothing)
 - 2. For each pixel (i,j) : (edge enhancement)
 - Compute the image gradient
 - $\nabla J(i,j) = (J_x(i,j), J_y(i,j))'$
 - Estimate edge strength
 - $e_s(i,j) = (J_x^2(i,j) + J_y^2(i,j))^{1/2}$
 - Estimate edge orientation
 - $e_o(i,j) = \arctan(J_x(i,j)/J_y(i,j))$
- The output are images E_s - Edge Strength - Magnitude
- and Edge Orientation E_o -


- E_s has large values at edges: Find local maxima
- ... but it also may have wide ridges around the local maxima (large values *around* the edges)



NONMAX_SUPPRESSION

- The inputs are E_s & E_o (outputs of CANNY_ENHANCER)
- Consider 4 directions $D=\{0,45,90,135\}$ wrt x
- For each pixel (i,j) do:
 1. Find the direction $d \in D$ s.t. $d \cong E_o(i,j)$ (normal to the edge)
 2. If $\{E_s(i,j)$ is smaller than at least one of its neigh. along $d\}$
 - $I_N(i,j)=0$
 - Otherwise, $I_N(i,j)=E_s(i,j)$
- The output is the thinned edge image I_N


Graphical Interpretation



Thresholding

- Edges are found by thresholding the output of NONMAX_SUPPRESSION
- If the threshold is too high:
 - Very few (none) edges
 - High MISDETECTIONS, many gaps
- If the threshold is too low:
 - Too many (all pixels) edges
 - High FALSE POSITIVES, many extra edges

SOLUTION: Hysteresis Thresholding




Canny Edge Detection (Example)

Original
image



gap is gone



Strong +
connected
weak edges



Strong
edges
only




Weak
edges




courtesy of G. Loy

Filters are templates

- Applying a filter at some point can be seen as taking a dot-product between the image and some vector
- Filtering the image is a set of dot products
- Insight
 - filters look like the effects they are intended to find
 - filters find effects they look like



Computer Vision - A Modern
Approach
Set: Linear Filters
Slides by D.A. Forsyth



Robinson Compass Masks

-1	0	1
-2	0	2
-1	0	1



0	1	2
-1	0	1
-2	-1	0



1	2	1
0	0	0
-1	-2	-1



2	1	0
1	0	-1
0	-1	-2



1	0	-1
2	0	-2
1	1	-1




0	-1	-2
-1	0	-1
2	1	0




-1	-2	-1
0	0	0
1	2	1



-2	-1	0
-1	0	1
0	1	2



Filter Bank



Leung & Malik, Representing and Recognizing the Visual
Apperance using 3D Textons, IJCV 2001